

# Frailty Classifier Project Modeling Strategy

Andrew Crane-Droesch

February 18, 2020

# Outline

1. Brief overview of the project
2. Annotation: why/what/how
3. Word embeddings
4. Modeling
5. Active learning

## Project overview

“Classify potentially preventable phenotypes of hospital admission among community-dwelling patients with chronic lung diseases using NLP within EHR data”

- ▶ Many hospital admissions are “low value”
  - ▶ Perhaps they could have been prevented
  - ▶ Patients often prefer not to be hospitalized
- ▶ It is not obvious *a priori* which people would be candidates for preventative measures, or what preventative measures to take
- ▶ This project builds on a related project that *identifies phenotypes* of people who are at high risk of hospitalization (“frailty”)
- ▶ The **immediate aim** is to identify *actionable aspects* of frailty
- ▶ Later, the goal is to use this to identify patients at high risk of being hospitalized, who have **actionable** risk reduction avenues
  - ▶ Pulmonary rehab might prevent some respiratory crises
  - ▶ Physical therapy might enable pulmonary rehab
  - ▶ Installing grab bars in bathrooms might prevent falls

# Project workflow

1. Download and process free text notes →
2. Annotation of spans of notes that are **positive** or **negative** for the following actionable aspects of frailty:
  - ▶ Fall risk
  - ▶ Musculo-skeletal problems
  - ▶ Respiratory impairment
  - ▶ Nutrition problems
3. Fitting models to predict the annotations from the spans
4. Using the fitted models to serve new batches of notes to annotate
5. Repeating 2, 3, and 4 until the models are good enough to classify text in the wild

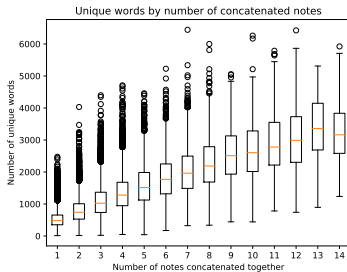
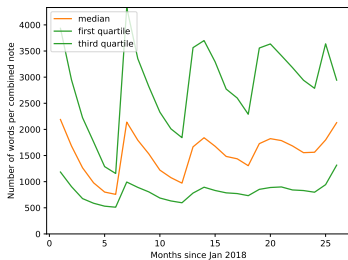
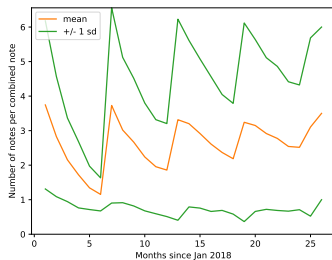
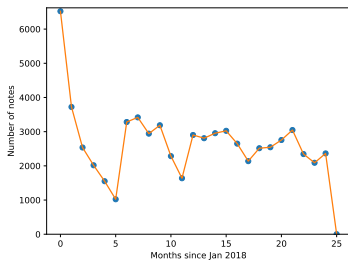
# Data munging and text processing

**Population:** all patients who've been at Penn for  $>1$  year, with at least one of several diagnoses relating to chronic lung disease.

Pipeline:

- ▶ Pull all signed **outpatient** progress notes from patients with relevant diagnoses
- ▶ Concatenate notes for individual patients into 6-month windows, simulating a prospective trial where targeting is done based on 6 months of history

# Summary stats



# Annotation

## Annotation

- 1 ----- [REDACTED] -----
- 2 [REDACTED] is\_a 75 y.o. male. he\_has a\_history\_of nodular\_sclerosing hodgkin\_lymphoma\_with pulmonary\_parenchymal involvement and sp [resp\_imp\_yes]
- 3 [resp\_imp\_yes]  
LLL endobronchial obstruction/RML extrinsic compression now resolved, hypothyroidism, prior 20py smoking\_history (quit in\_his 40s),\_a  
he\_was\_diagnosed with lymphoma in [REDACTED] based\_on supraclavicular\_lymph\_node biopsy and treated\_with AVD with last treatr  
showed\_gradual\_improvement with\_complete metabolic\_response on\_june scan, followed\_by increasing fdg\_activity in abdominal\_lymp  
also\_demonstrated new, numerous, small perilymphatic nodules concerning\_for recurrent lymphoma vs sarcoid-like\_reaction.
- 4 Subsequent bronchoscopy\_confirmed sarcoid. he\_did\_not tolerate prednisone due\_to irritability\_and was hesitant\_to try other\_immunos  
imaging at\_this\_time.
- 5 I last saw\_him on [REDACTED] at\_that\_time, I recommended: - Stop bactrim, famotidine - PFT today - PFT, 6MWT at next\_visit in 3 mo  
develops new respiratory, cardiac, dermatologic, or\_ocular symptoms - will\_perform EKG at next\_visit respiratory\_symptoms\_are stable
- 6 [resp\_imp\_yes]  
Dyspnea after hot shower, prolonged moderate\_exertion such\_as walking outside home for\_more\_than 10 minutes.
- 7 Has an albuterol nebulizer at\_home\_but has\_not tried using this to\_see\_if it alleviates symptoms.
- 8 lymphoma\_is stable.
- 9 [msk\_prob\_yes]  
Planning to\_have hip\_replacement in Fall as joint\_pain often limits his activity.
- 10 [msk\_prob\_yes]  
Will do PT afterwards.

# Tokenization and label processing

```
>>> i+=20; x.iloc[i:(i+20)]
```

	token	length	end	start	Frailty_nos	Msk_prob	Nutrition	Resp_imp	Fall_risk
560	,	2	3190	3188	0	0	-1	0	0
561	well	4	3194	3190	0	0	-1	0	0
562	-	1	3195	3194	0	0	-1	0	0
563	nourished	9	3204	3195	0	0	-1	0	0
564	,	2	3206	3204	0	0	0	0	0
565	and	4	3210	3206	0	0	0	0	0
566	in	3	3213	3210	0	0	0	0	0
567	no	3	3216	3213	0	0	0	0	0
568	distress	8	3224	3216	0	0	0	0	0
569	.	2	3226	3224	0	0	0	0	0
570	No	3	3229	3226	0	0	0	0	0
571	distress	8	3237	3229	0	0	0	0	0
572	.	1	3238	3237	0	0	0	0	0
573	\n	2	3240	3238	0	0	0	0	0
574	HENT	4	3244	3240	0	0	0	0	0
575	:	1	3245	3244	0	0	0	0	0
576	\n	2	3247	3245	0	0	0	0	0
577	Head	4	3251	3247	0	0	0	0	0
578	:	2	3253	3251	0	0	0	0	0
579	Normocephalic	14	3267	3253	0	0	0	0	0

```
>>>
```

↑ this is the raw data. How can we use the token to predict the label?



Word embeddings

# Document matrix

Start with a sparse matrix “one-hot” encoded representation of a document:

- ▶ The rows correspond to positions in a sequence
- ▶ The columns correspond to unique words
- ▶ The sum of the matrix equals the number of rows

$D \equiv$		<i>I</i>	<i>am</i>	<i>Sam</i>	<i>would</i>	<i>you</i>	<i>like</i>
	<i>I</i>	1	0	0	0	0	0
	<i>am</i>	0	1	0	0	0	0
	<i>Sam</i>	0	0	1	0	0	0
	<i>Sam</i>	0	0	1	0	0	0
	<i>I</i>	1	0	0	0	0	0
	<i>am</i>	0	1	0	0	0	0
	<i>would</i>	0	0	0	1	0	0
	<i>you</i>	0	0	0	0	1	0
	<i>like</i>	0	0	0	0	0	1

Call it  $D$ . It has dimensions  $N \times V$  – total words by unique words.

# Document matrix

- ▶ Document matrices are *sparse*. Lots of zeros.
- ▶ They are high-dimensional. Using  $\mathbf{D}$  directly as a design matrix in a model is generally inefficient.
- ▶ A document matrix doesn't explicitly encode any information about how words might be similar to each other. Synonyms are wholly different from each other.
  - ▶ The euclidean distance between any two word vectors (rows in the matrix) will always be  $\sqrt{2}$
  - ▶ The cosine similarity will always be zero
- ▶ Summing them column-wise makes a unigram matrix:

	$I$	$am$	$Sam$	$would$	$you$	$like$
$1 - gram$	2	2	2	1	1	1

# Embeddings

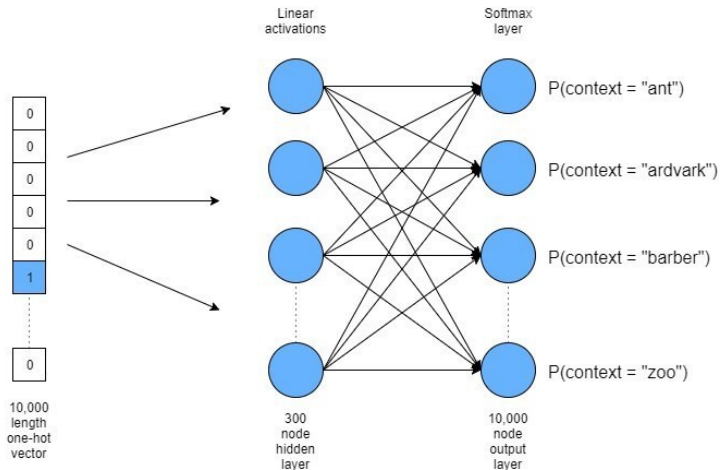
Embeddings attempt to do the following:

- ▶ Reduce the dimensionality of  $\mathbf{D}$  without losing too much information
- ▶ Capture the similarity between similar words

Different embedding methods do this in different ways, but all create analogous output:


- ▶ word2vec, fasttext, GloVe

# Embeddings in pictures



# Embeddings in pictures

## Cosine Similarity



# Embeddings in math

(a simple) Embedding model:

$$\mathbf{D} = s(\mathbf{D}\mathbf{\Gamma}_0\mathbf{\Gamma}_1) + \epsilon$$

where:

- ▶  $\mathbf{D}$  is the document matrix, dimension  $N \times V$
- ▶  $\mathbf{\Gamma}_0$  is the first weight matrix. It has dimension  $V \times U$ , where  $U \ll V$ . It is responsible for “embedding”  $\mathbf{D}$  in a lower-dimensional space
- ▶  $\mathbf{\Gamma}_1$  is the second weight matrix. Its dimension is  $U \times V$ . It is responsible for taking the lower-dimensional representation and bringing it back to the original resolution.
- ▶  $s()$  is the softmax function. It maps the maps real numbers to a vector of probabilities summing to 1.
- ▶  $\epsilon$  is the error to be minimized. As the bottleneck gets skinnier (i.e.: as  $U$  gets smaller),  $\epsilon$  will have to increase.

# Embeddings in math

(a simple) Embedding model:

$$\mathbf{D} = s(\mathbf{D}\mathbf{\Gamma}_0\mathbf{\Gamma}_1) + \epsilon$$

where:

- ▶ The embedded text, call it  $\mathbf{E}$ , is defined as  $\mathbf{D}\mathbf{\Gamma}_0$  (dimension  $N \times U$ )
- ▶ (most) Embedding models are trained by computing the derivatives of the parameter matrices  $\mathbf{\Gamma}_0, \mathbf{\Gamma}_1$  with respect to a loss function, changing the parameters to make the loss smaller, and then recomputing the derivatives and repeating. This is backpropagation.
- ▶ the error  $\epsilon$  is implicit. The goal is to minimize it, to get the best possible  $\mathbf{\Gamma}$ . But as the dimension of  $\mathbf{\Gamma}$  gets smaller, this gets harder to do.



## Embeddings in practice

Training and using embeddings are separate processes. Trained embedding models give you  $\mathbf{\Gamma}_0$ . You then supply your own text to multiply it against. For example, take  $\mathbf{D}$  to be our Dr Seuss example from before. The embeddings  $\mathbf{\Gamma}_0$  might look like this:

	$u1$	$u2$
$I$	1	2
$am$	3	4
$Sam$	5	6
$would$	7	8
$you$	9	10
$like$	11	12

(Note that the entries of the matrix are totally unrealistic in this example)

## Embeddings in practice

Say you wanted to embed the phrase “I like Sam” into this two-dimensional space:

$$\left[ \begin{array}{c|cccccc} & & & & & & \\ I & 1 & 0 & 0 & 0 & 0 & 0 \\ like & 0 & 0 & 0 & 0 & 0 & 1 \\ Sam & 0 & 0 & 1 & 0 & 0 & 0 \end{array} \right] \left[ \begin{array}{c|cc} & u1 & u2 \\ \hline I & 1 & 2 \\ am & 3 & 4 \\ Sam & 5 & 6 \\ would & 7 & 8 \\ you & 9 & 10 \\ like & 11 & 12 \end{array} \right] = \left[ \begin{array}{c|cc} & u1 & u2 \\ \hline I & 1 & 2 \\ like & 11 & 12 \\ Sam & 5 & 6 \end{array} \right]$$

Because the the document matrix  $\mathbf{D}$  has only one “1” per row, this matrix multiplication is the same as a “lookup” operation.

Still, the matrix notation will come in handy later.

# Word embeddings to sentence or span embeddings

Supervised learning jobs can all be conceptualized as

$$y = f(\mathbf{X}) + \epsilon$$

where  $y$  is some label and  $\mathbf{X}$  is a matrix of features. So let's add a label to our toy example:

	<i>I</i>	<i>am</i>	<i>Sam</i>	<i>would</i>	<i>you</i>	<i>like</i>	<b>label</b>
<i>I</i>	1	0	0	0	0	0	0
<i>am</i>	0	1	0	0	0	0	0
<i>Sam</i>	0	0	1	0	0	0	1
<i>Sam</i>	0	0	1	0	0	0	0
<i>I</i>	1	0	0	0	0	0	0
<i>am</i>	0	1	0	0	0	0	0
<i>would</i>	0	0	0	1	0	0	1
<i>you</i>	0	0	0	0	1	0	1
<i>like</i>	0	0	0	0	0	1	0

(I labeled words that come after a verb)

# Word embeddings to sentence or span embeddings

If we fit

$$\mathbf{label} = f(\mathbf{D}) + \epsilon$$

or

$$\mathbf{label} = f(\mathbf{E}) + \epsilon$$

our model would do nothing but capture the correlations between individual words and labels. But what if we want to capture the surrounding context?

First, let's remember that the matrix equation above is the same as

$$label_i = f(E_i) + \epsilon$$

where  $i$  indexes rows from 1 to  $N$ . A natural strategy could be to include lags and leads:

$$label_i = f(E_i, E_{i-1}, E_{i+1}) + \epsilon$$

But lots of other possibilities exist.

# Word embeddings to sentence or span embeddings

Consider the span “I am Sam”. Its embedding is:

	<i>u1</i>	<i>u2</i>	<b>label</b>
<i>I</i>	1	2	0
<i>am</i>	3	4	0
<i>Sam</i>	5	6	1

The centroid word is “am”, which is associated with a label of zero. How do we associate information about the surrounding words with that zero label? Several options, all of which involve “windowing”:

- ▶ Lags and leads
- ▶ span-wise maxes and mins, averages

	<i>u1</i>	<i>u2</i>	<b>label</b>	<i>lag</i>	<i>lead</i>	<i>min</i>	<i>max</i>
<i>am</i>	3	4	0	1 2	5 6	1 2	5 6

In this toy example, the lags and leads happen to be the same as the maxes and mins, but that won't generally be the case.

## Word embeddings to sentence or span embeddings

We've been working with a bandwidth of 1 (as measured from the centroid.). This is the same thing as a window size of 3.

Taking an average over this window is the same as doing the following:

$$[.33, .33, .33] \left[ \begin{array}{c|cc} & u1 & u2 \\ \hline I & 1 & 2 \\ am & 3 & 4 \\ Sam & 5 & 6 \end{array} \right] = [3, 4]$$

But why weight evenly? We could do

$$[.2, .7, .1] \left[ \begin{array}{c|cc} & u1 & u2 \\ \hline I & 1 & 2 \\ am & 3 & 4 \\ Sam & 5 & 6 \end{array} \right] = [2.8, 3.8]$$

Optimal weighting schemes could be hyperparameters. Or they could be *directly estimated* at the same time as model parameters if we're using a neural network. This is basically the same thing as a one-dimensional *convolution*.

# Embedded Ngrams

Ngrams are usually computed over the whole of a document, but they could also be computed over a window. Here's bigrams for our dummy example:

		$(I, I)$	$(I, am)$	$(I, sam)$	$\dots$	$(am, Sam)$	$(Sam, Sam)$
$G^2 =$	$I$	0	1	0	$\dots$	0	0
	$am$	0	1	0	$\dots$	1	0
	$Sam$	0	0	0	$\dots$	1	0

These bigrams are going to be much too high-dimensional to use as-is, but we can embed them into a lower-dimensional space:

$$E^2 = \underbrace{G^2}_{N \times V^2} \left[ \underbrace{\Gamma_0 \otimes \Gamma_0}_{V^2 \times U^2} \right]$$

Doing so could help capture topology-dependent phenomena like negation and adjectives/adverbs. The dimensionality is too high to do this with the whole vocab even at  $U = 100$ , but this could be reduced through TF-IDF weighting of bigrams and PCA on the matrix  $\Gamma$ .

# Models

Given a label and a transformation that takes the document matrix  $D$  and returns  $X$ , fitting a statistical model is generally a fairly straightforward exercise of tuning hyperparameters.

Neural nets are the exception, because they effectively create their own design matrices through representation learning. I'll discuss those later.

We'll try:

- ▶ Penalized logistic regression
- ▶ Random forest
- ▶ XGboost (with trees, probably)



# Structured data (not incorporated yet)

We'll be using:

- ▶ Labs
- ▶ Comorbidities
- ▶ Frailty indices from the literature
- ▶ Vitals
- ▶ Demographics

All of the structure data will only vary at the *patient level*. We'll have  $>> 10^6$  labeled tokens, but from  $<< 1000$  labeled patients.

# Random Forest

Doing the RF as a first pass, because RF's require minimal tuning.

Key considerations:

- ▶ The data aren't IID:  $X_i$  will be deeply correlated with  $X_{i+1}$ 
    - ▶ because the features include window statistics
    - ▶ because the tokens represent language, which is non-random
- Training and test sets need to be *cluster-sampled*
- ▶ Labels are rare, and classification forests aren't probabilistically calibrated: their score isn't in general the expectation of a class's probability
- The **ranger** package (in R) offers the *probability forest* option. Terminal nodes output probabilities, rather than majority votes. This will help in the early stages of active learning.


# Featurization

Here are the sets of features of the embeddings that we're using:


- ▶ Identity:  $f(i) = E_i$
- ▶ Last word (one lag):  $f(i) = E_{i-1}$
- ▶ Last word (two lags):  $f(i) = E_{i-2}$
- ▶ Weighted mean over bandwidth ( $b$ ):  $f(i) = kE_{i \in b}$ 
  - ▶ where  $k$  is a  $1 \times b$  vector of interpolated densities from under a standard normal between -3 and 3
- ▶ Max over bandwidth:  $f(i) = \text{lambda } x: \text{ np.amax}(x, \text{ axis}=0)$
- ▶ Min over bandwidth:  $f(i) = \text{lambda } x: \text{ np.amin}(x, \text{ axis}=0)$

Bandwidth set to 30


# Results (preliminary)



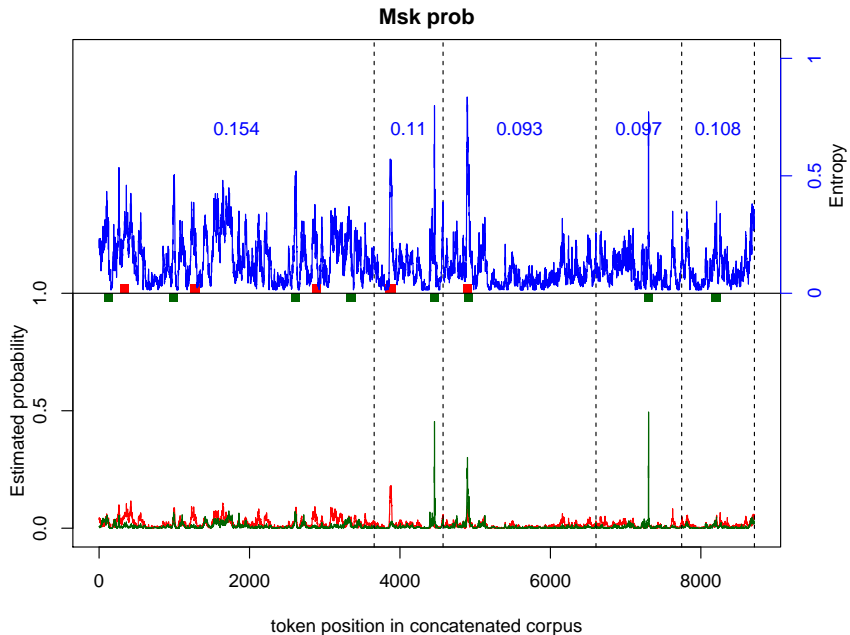
# Results (preliminary)



# Results (preliminary)



# Results (preliminary)



# Neural nets

- ▶ I haven't trained a neural net yet
- ▶ I expect that it will be much more efficient
- ▶ Neural nets have the nice feature that they implicitly do the featurization for you
  - ▶ That's what *representation learning* is, basically
- ▶ That doesn't mean that featurizations fed to the other models wouldn't be useful to the net
  - ▶ Basically it'd be giving the AI a cribsheet
- ▶ Likewise it'll be important to think carefully about model architecture, especially around questions of context (negation, allusion, etc.)



# Active learning

How do we take the output of models fit to the first 25 patient notes and use them to pick the most fruitful notes to annotate next?

- ▶ Standard approach is to pick examples that we're maximally uncertain about
- ▶ *Entropy* is a useful measure of uncertainty from information theory

$$\text{Entropy} \equiv H = - \sum_c p(\text{class} = c) \log(p(\text{class} = c))$$

- ▶ It is maximized when you know nothing about which class an example belongs to.
- ▶ We're going to pick notes that have a lot of entropy.
- ▶ Not obvious how best to aggregate entropy over the note. Thinking perhaps an average over the top quantile. Maybe decile.

# Active learning

Example: In the beginning, the model isn't very good and will default to the dominant class (neutral)

- ▶ A token that the model thinks is neutral will have low entropy:  $H([.01, .98, .01] = .11)$
- ▶ If the model thinks that *maybe* another token might be positive, entropy will be a little lower:  $H([.1, .88, .01] = .39)$ . That token will make the note more likely to get picked.

Eventually, the model will get better about distinguishing the positive classes:

- ▶ After a while, it gets confident about a negative:  $E([.01, .01, .98] = .11)$ , and that token won't influence what notes get picked to annotate
- ▶ If a note has examples that we are completely uninformed about (maximum entropy), we'd gravitate towards those notes:  $H([.33, .33, .33] = 1.1)$

# Meta-considerations

This isn't a causal model, so the usual caveats about feature importance apply

- ▶ That said, it'd be useful to know what sort of features trigger positives and negatives for these frailty aspects
- ▶ Standard permutation importance is usually only done for RF's but can be coded for any model

Fairness:

- ▶ What happens if our model is much less certain about patients who are/aren't black/white/men/women/rich/indigent?
- ▶ Do we think it's likely that clinicians write differently about patients who are different – consciously or unconsciously, intentionally or otherwise?
- ▶ Can we apply post-hoc statistical fixes to correct for such biases?

# Wrapup

Takeaways, challenges, and points to discuss:

- ▶ We are annotating and predicting labels on spans of text, but
  - ▶ how do we classify actual patients?
- ▶ Context is critical in text modeling.
  - ▶ What are the best ways to represent it?
- ▶ After the first batches of notes, we're going to pick subsequent notes about which we're maximally uncertain.
  - ▶ But what is the best way to aggregate that uncertainty?
- ▶ Bias is likely to pop up along the way.
  - ▶ How can we plan for it?

**Thanks!**